





Cold
Spring
Harbor
Laboratory

RNA-Seq Module 2: **FASTQ/SAM/BAM/BED file formats**

Arpad Danos, Felicia Gomez, Obi Griffith, Malachi Griffith,
My Hoang, Mariam Khanfar, Chris Miller, Kartik Singhal

Advanced Sequencing Technologies & Bioinformatics Analysis November 10-23, 2024



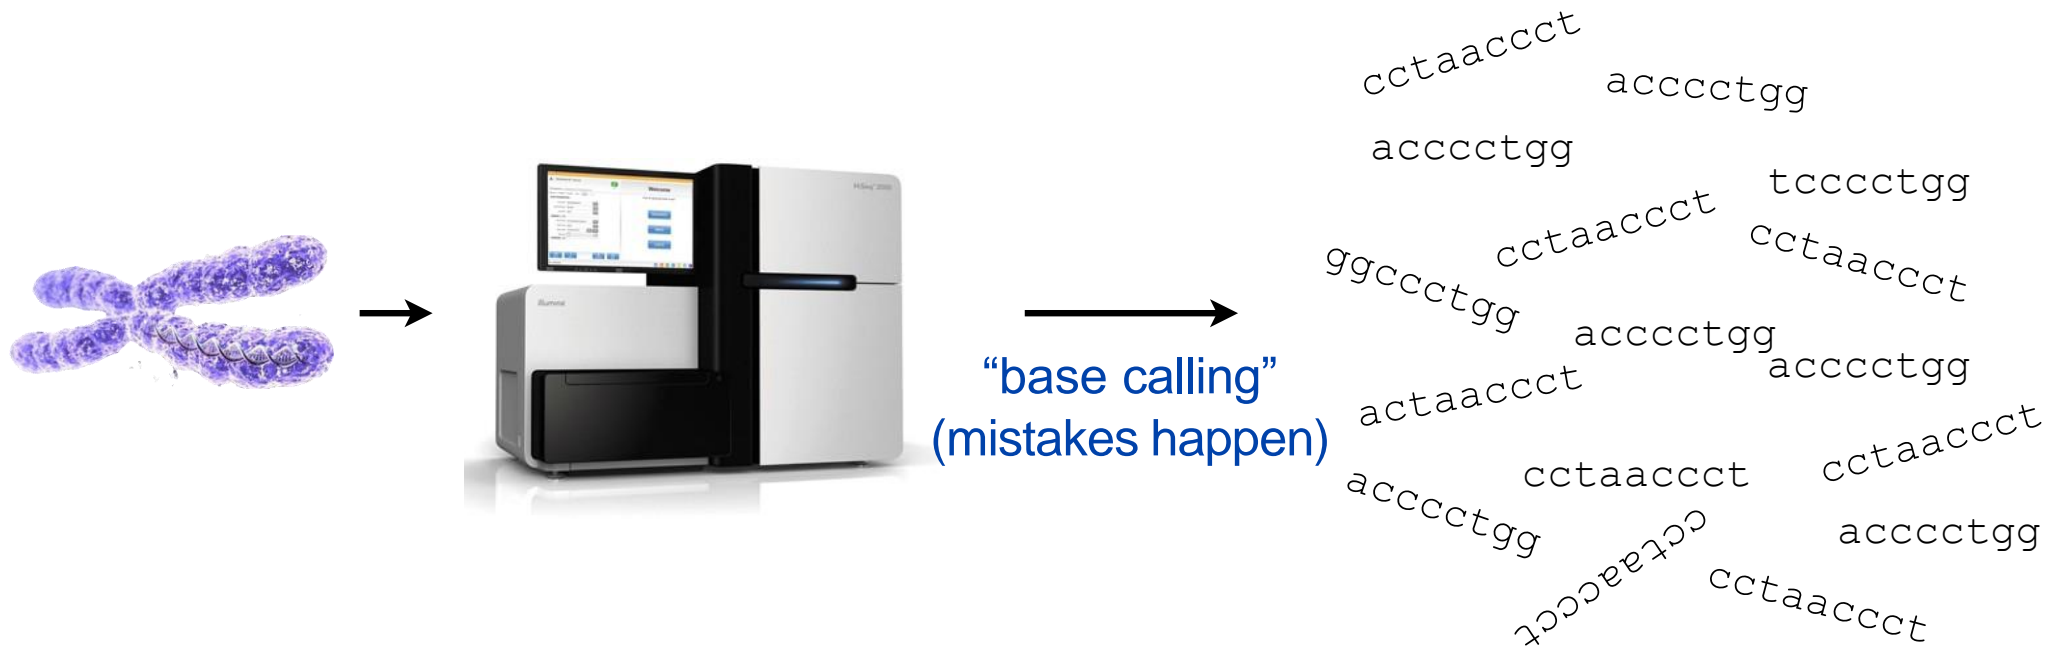
Washington University in St. Louis
SCHOOL OF MEDICINE

Outline

- What is a sequence read?
- Unaligned read
- Aligned read
- Difference between the stored aligned reads file formats

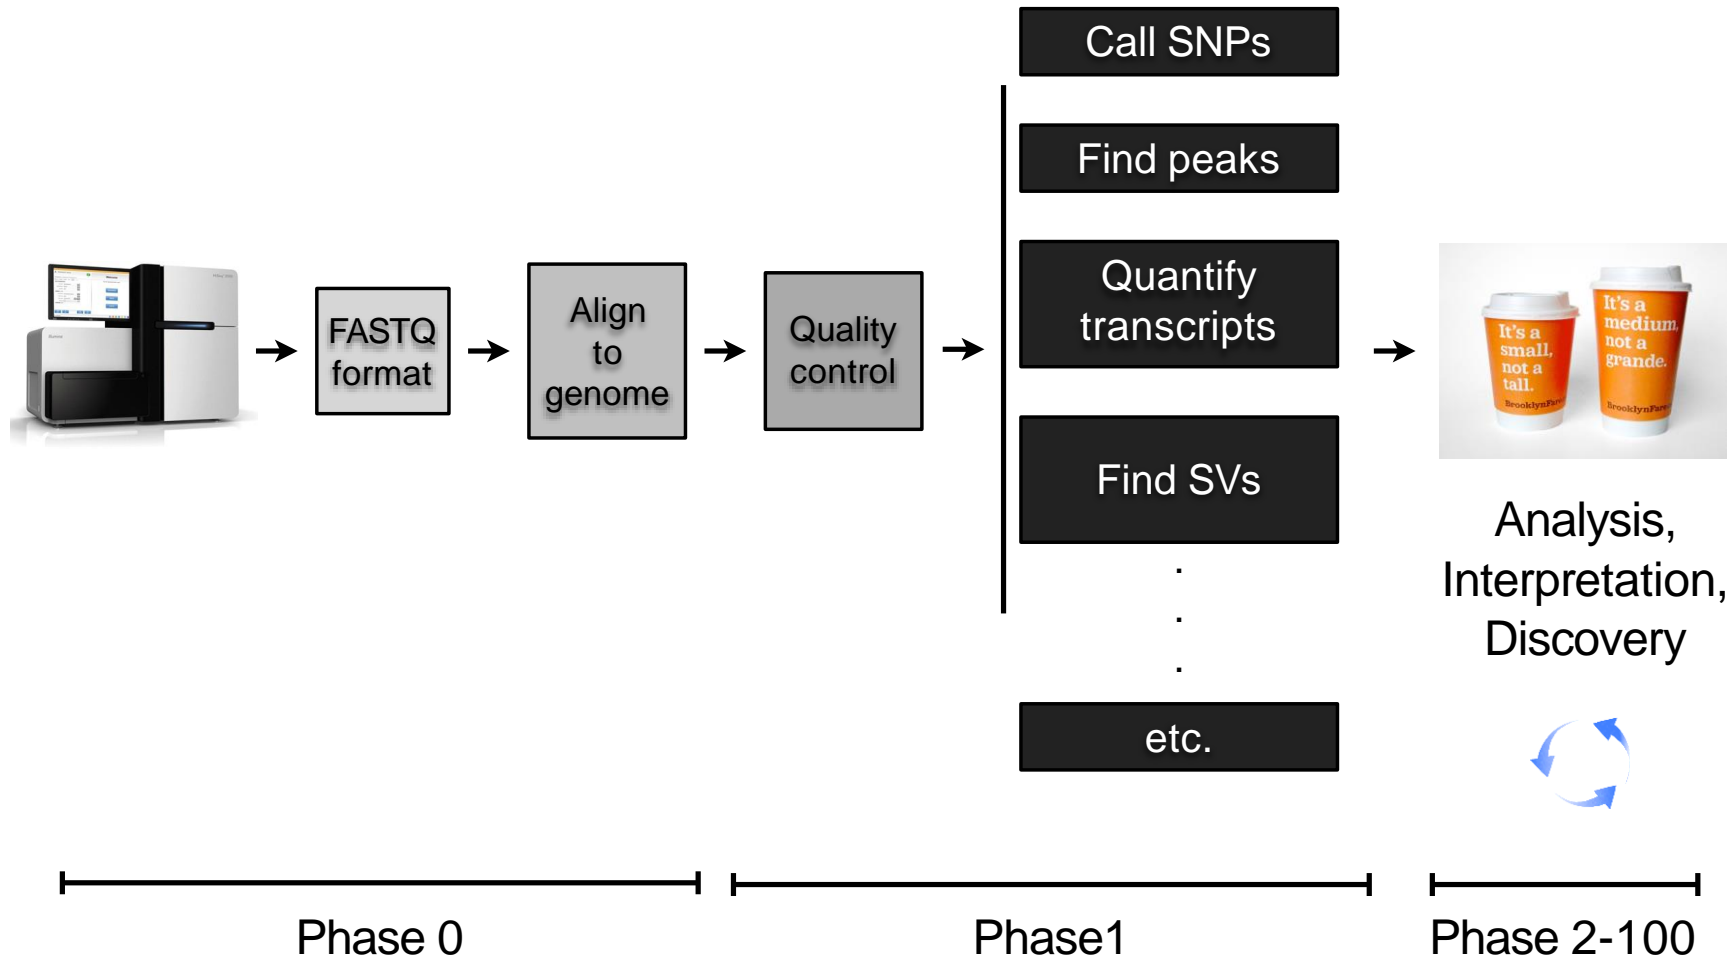
Reads are the sequencer's best guess at what it saw for a given DNA molecule.

It's the "raw" data.



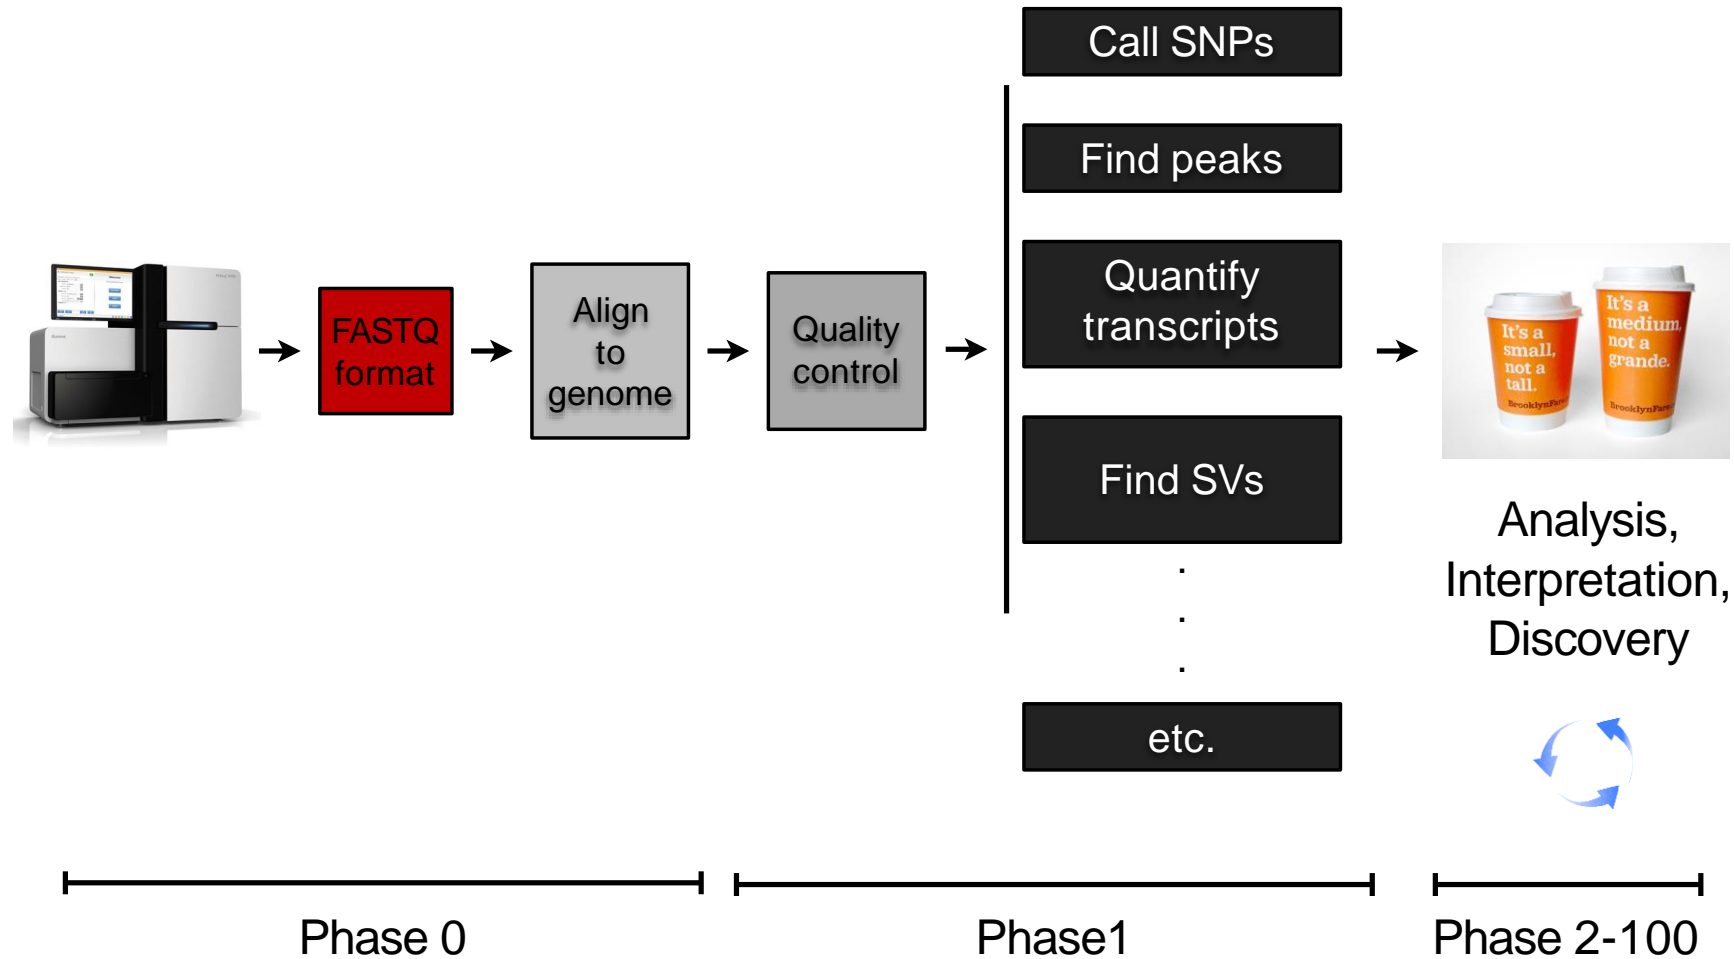
Slide courtesy of Andrew Farrell

Alignment is central to most genomics applications



Slide courtesy of Andrew Farrell

Alignment is central to most genomics applications



Slide courtesy of Andrew Farrell

The FASTQ format

A “standard” format for storing and defining sequences from next-generation sequencing technologies.

Sequence ID	-----●	@SEQ_ID
Sequence	—————●	GATTTGGGGTTCAAAGCAGTATCGATCAAATAGTAAATCCATTTGTTCAACTCACAGTTT
<separator>	—————●	+
Quality scores	—————●	! ' ' * ((((* * * +)) % % % + +) (% % % %) . 1 * * * - + * ' ')) * * 5 5 C C F > > > > > C C C C C C C C 6 5

- FASTQ files are generally used to store short-read data from high-throughput sequencing experiments
- The sequence and quality scores are usually put into a single line

Slide courtesy of Andrew Farrell

Sequence IDs

@HWUSI-EAS100R:6:73:941:1973#0/1

HWUSI-EAS100R	the unique instrument name
6	flowcell lane
73	tile number within the flowcell lane
941	'x'-coordinate of the cluster within the tile
1973	'y'-coordinate of the cluster within the tile
#0	index number for a multiplexed sample (0 for no indexing)
/1	the member of a pair, /1 or /2 (<i>paired-end or mate-pair reads only</i>)

@EAS139:136:FC706VJ:2:2104:15343:197393 1:Y:18:ATCACG

EAS139	the unique instrument name
136	the run id
FC706VJ	the flowcell id
2	flowcell lane
2104	tile number within the flowcell lane
15343	'x'-coordinate of the cluster within the tile
197393	'y'-coordinate of the cluster within the tile
1	the member of a pair, 1 or 2 (<i>paired-end or mate-pair reads only</i>)
Y	Y if the read is filtered, N otherwise
18	0 when none of the control bits are on, otherwise it is an even number
ATCACG	index sequence

Slide courtesy of Andrew Farrell

Quality scores

Sequence ID•	@SEQ_ID
Sequence•	GATTTGGGGTTCAAAGCAGTATCGATCAAATAGTAAATCCATTTGTTCAACTCACAGTTT
<separator>•	+
Quality scores•	! ' * (((* * * +)) % % % + +) (% % % %) . 1 * * * - + * ' ') * * 5 5 C C F > > > > > C C C C C C C 6 5



Example:

ASCII characters: Code →

ASCII characters: Value →

!	+	E	I
↙	↓	↓	↘
33	43	69	73

Quality scores

Sequence ID•	@SEQ_ID
Sequence•	GATTTGGGGTTCAAAGCAGTATCGATCAAATAGTAAATCCATTTGTTCAACTCACAGTTT
<separator>•	+
Quality scores•	! ' * (((* * * +)) % % % + +) (% % % %) . 1 * * * - + * ' ') * * 5 5 C C F > > > > > C C C C C C C 6 5



Example:

ASCII characters: Code →

! + E I

ASCII characters: Value →

33 43 69 73

ASCII is inferred by the probability of error ~ Q score (Phred score)

Quality scores

Sequence ID●	@SEQ_ID
Sequence●	GATTTGGGGTTCAAAGCAGTATCGATCAAATAGTAAATCCATTTGTTCAACTCACAGTTT
<separator>●	+
Quality scores●	! '* ((((* * * +)) % % % + +) (% % % %) . 1 * * * - + * ' ') * * 5 5 C C F > > > > > C C C C C C C 6 5



Qualities are based on the Phred scale and are *encoded*

$$Q = -10 \cdot \log_{10}(P_{\text{err}})$$

$$Q = -10 \cdot \log_{10}(0.01)$$

$$Q = 20 \text{ (Q is the Phred score)}$$

- FASTQ files encodes phred scores as ASCII characters
- Phred quality scores characterize the quality of DNA sequences – these scores are assigned by the sequencer
- A quality score of 20 (Q20) represents an error rate of 1 in 100 (meaning every 100 bp sequencing read may contain an error); call accuracy of 99%

Slide courtesy of Andrew Farrell

Quality score encoding

Dec	Hex	Char	Dec	Hex	Char	Dec	Hex	Char	Dec	Hex	Char
0	00	Null	32	20	Space	64	40	@	96	60	`
1	01	Start of heading	33	21	!	65	41	A	97	61	a
2	02	Start of text	34	22	"	66	42	B	98	62	b
3	03	End of text	35	23	#	67	43	C	99	63	c
4	04	End of transmit	36	24	\$	68	44	D	100	64	d
5	05	Enquiry	37	25	%	69	45	E	101	65	e
6	06	Acknowledge	38	26	&	70	46	F	102	66	f
7	07	Audible bell	39	27	'	71	47	G	103	67	g
8	08	Backspace	40	28	(72	48	H	104	68	h
9	09	Horizontal tab	41	29)	73	49	I	105	69	i
10	0A	Line feed	42	2A	*	74	4A	J	106	6A	j
11	0B	Vertical tab	43	2B	+	75	4B	K	107	6B	k
12	0C	Form feed	44	2C	,	76	4C	L	108	6C	l
13	0D	Carriage return	45	2D	-	77	4D	M	109	6D	m
14	0E	Shift out	46	2E	.	78	4E	N	110	6E	n
15	0F	Shift in	47	2F	/	79	4F	O	111	6F	o
16	10	Data link escape	48	30	0	80	50	P	112	70	p
17	11	Device control 1	49	31	1	81	51	Q	113	71	q
18	12	Device control 2	50	32	2	82	52	R	114	72	r
19	13	Device control 3	51	33	3	83	53	S	115	73	s
20	14	Device control 4	52	34	4	84	54	T	116	74	t
21	15	Neg. acknowledge	53	35	5	85	55	U	117	75	u
22	16	Synchronous idle	54	36	6	86	56	V	118	76	v
23	17	End trans. block	55	37	7	87	57	W	119	77	w
24	18	Cancel	56	38	8	88	58	X	120	78	x
25	19	End of medium	57	39	9	89	59	Y	121	79	y
26	1A	Substitution	58	3A	:	90	5A	Z	122	7A	z
27	1B	Escape	59	3B	;	91	5B	[123	7B	{
28	1C	File separator	60	3C	<	92	5C	\	124	7C	
29	1D	Group separator	61	3D	=	93	5D]	125	7D	}
30	1E	Record separator	62	3E	>	94	5E	^	126	7E	~
31	1F	Unit separator	63	3F	?	95	5F	_	127	7F	□

Formula for getting PHRED quality from encoded quality:

$$Q = \text{ascii(char)} - 33$$

Example:

! + E I

ASCII
-33 (Q)

33	43	69	73
0	10	36	40

- ASCII = **A**merican **S**tandard **C**ode for **I**nformation **I**nterchange
- Every text symbol must have an integer value representing it inside the computer
- An ASCII code is the numerical representation of a character such as 'a' or '@'

Image courtesy of Andrew Farrell

$$Q = -10 \log_{10}(P_{\text{error}})$$

Probability of Error		Q
1/1,000,000	0.000001	60
1/100,000	0.000010	50
1/10,000	0.000100	40
1/1,000	0.001000	30
1/100	0.010000	20
1/10	0.100000	10
1/1	1.000000	0

- **Higher Q scores** indicate a smaller probability of error.
- **Lower Q scores** indicate lower confidence in the called base.
 - Increased false-positive variant calls
- Q30 is a standard a benchmark for quality in next-generation sequencing

Slide courtesy of Andrew Farrell

FASTQ Report Summary

FastQC Report

Tue 12 Jan 2016
2009-08_lib324_miseq_r0030_251bp_R1.fastq.gz


Summary

- ✓ [Basic Statistics](#)
- ✗ [Per base sequence quality](#)
- ✗ [Per tile sequence quality](#)
- ✓ [Per sequence quality scores](#)
- ✗ [Per base sequence content](#)
- ✓ [Per sequence GC content](#)
- ✓ [Per base N content](#)
- ! [Sequence Length Distribution](#)
- ✓ [Sequence Duplication Levels](#)
- ✓ [Overrepresented sequences](#)
- ✓ [Adapter Content](#)
- ! [Kmer Content](#)

✓ Basic Statistics

Measure	Value
Filename	2009-08_lib324_miseq_r0030_251bp_R1.fastq.gz
File type	Conventional base calls
Encoding	Sanger / Illumina 1.9
Total Sequences	503810
Sequences flagged as poor quality	0
Sequence length	35-251
%GC	66

✗ Per base sequence quality

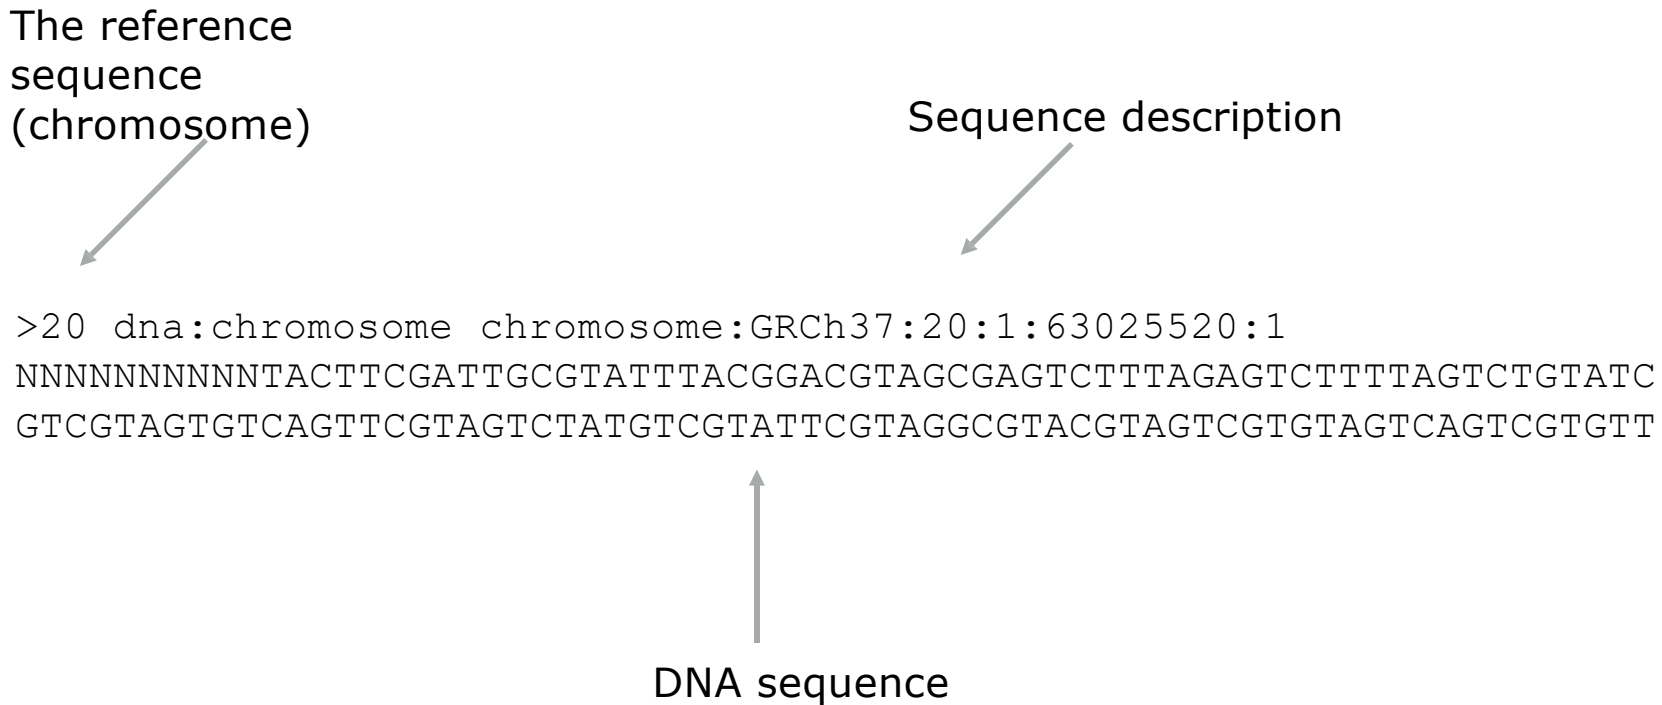


FASTA format

We start with a reference genome to map to

The reference
sequence
(chromosome)

Sequence description




```
>20 dna:chromosome chromosome:GRCh37:20:1:63025520:1
NNNNNNNNNTACTTCGATTGCGTATTTACGGACGTAGCGAGTCTTTAGAGTCTTTTAGTCTGTATC
GTCGTAGTGTCAGTTCGTAGTCTATGTCGTATTCGTAGGCGTACGTAGTCGTGTAGTCAGTCGTGTT
```

DNA sequence

http://en.wikipedia.org/wiki/FASTA_format

Slide courtesy of Andrew Farrell

Aligning to a reference genome; the crucial first step




Could fit here - but there
are differences

Could fit here as well.

Slide courtesy of Andrew Farrell

Single-end alignment



1. Sequence the **entire length** of each molecule the library

472bp 499bp
519bp
503bp 561bp

2. Align each **contiguous** molecule to the reference genome.

472bp 519bp

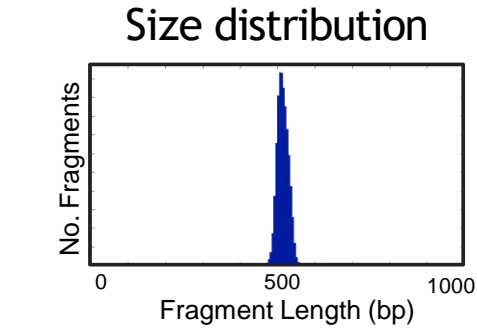
503bp

499bp

561bp

Slide courtesy of Andrew Farrell


Paired-end alignment




0. Size-selected
DNA library (e.g., ~500bp)



1. Sequence ***solely the 5' ends*** of
each molecule the library




Paired End sequencing




Slide courtesy of Andrew Farrell

Paired-end alignment



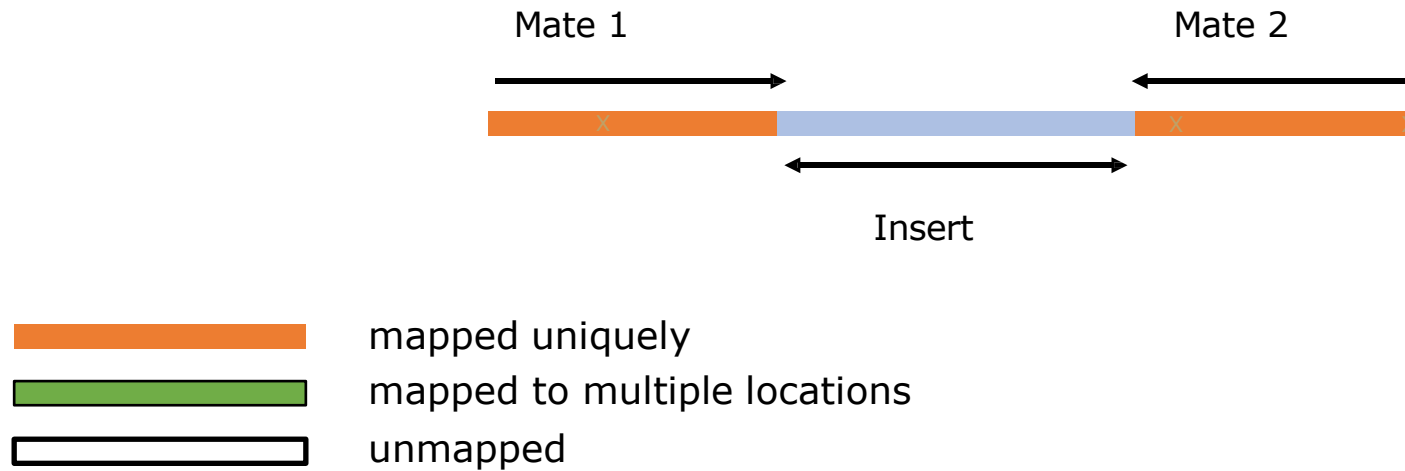
Slide courtesy of Andrew Farrell

Paired-end alignment



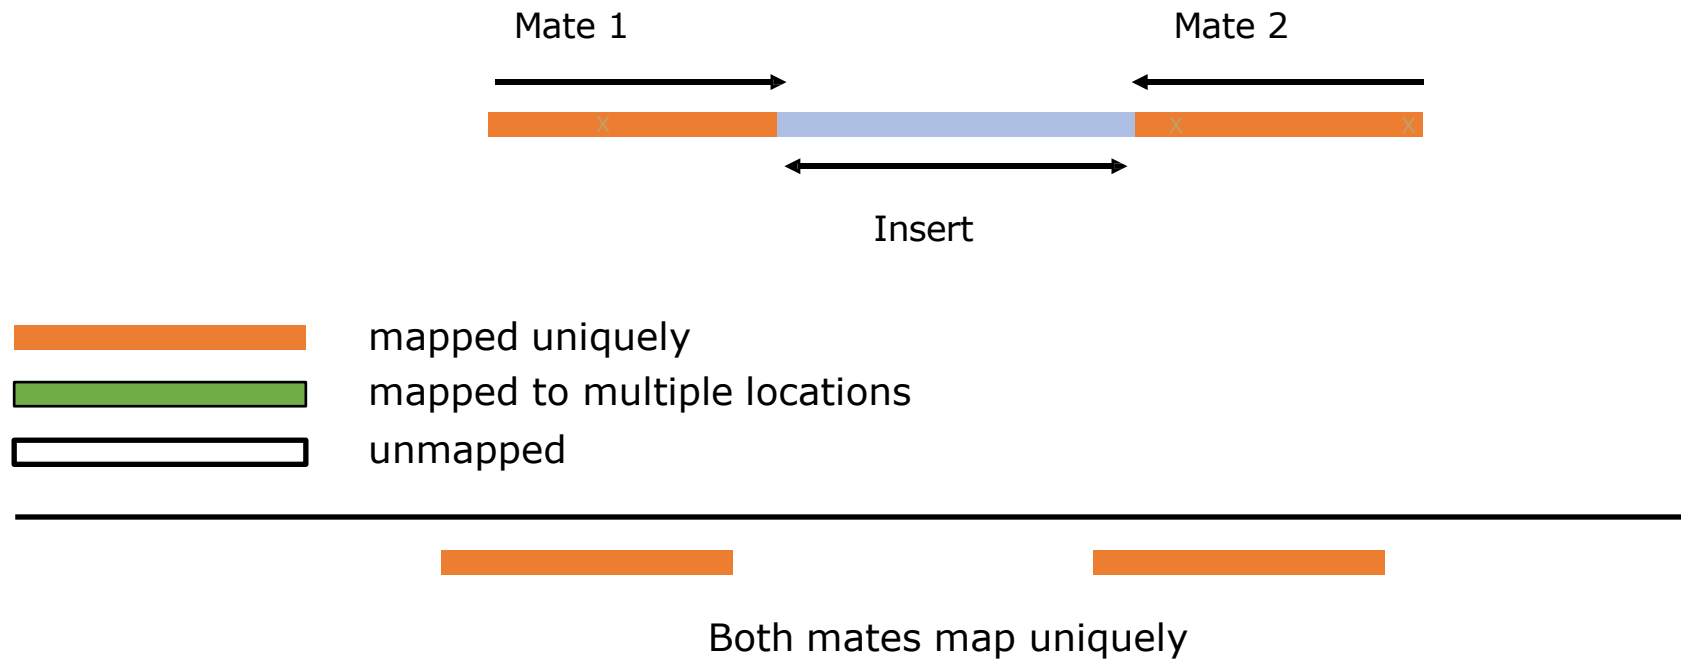
Slide courtesy of Andrew Farrell

Paired-end alignment



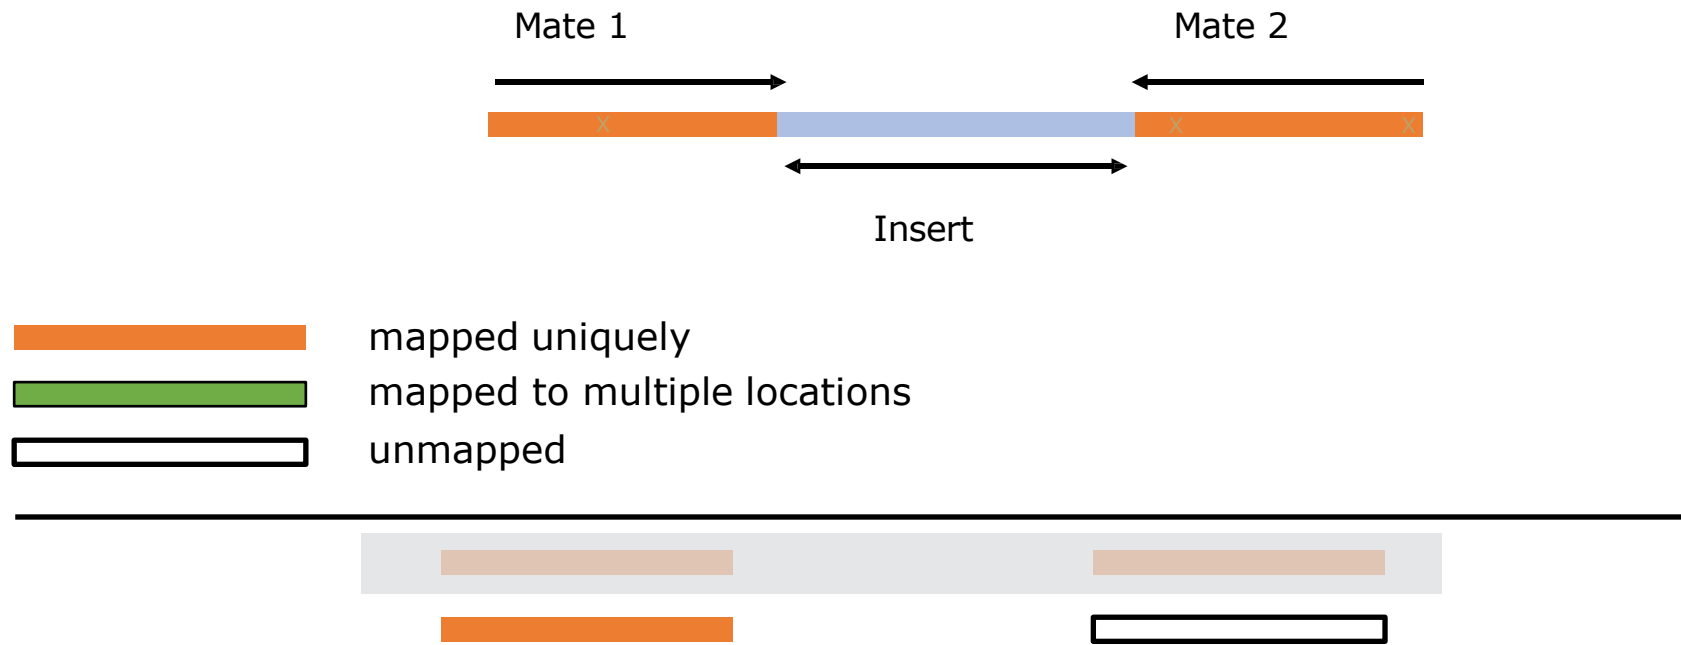
Slide courtesy of Andrew Farrell

Paired-end alignment




Slide courtesy of Andrew Farrell




Paired-end alignment




One mate maps uniquely, the other is unmapped

Paired-end alignment




-  mapped uniquely
-  mapped to multiple locations
-  unmapped




One mate maps uniquely, the other maps to multiple locations

Use fragment length distribution to determine most likely location



What needs to be stored?



Where did the read map?

How confident are we that we are correct?

Which strand does the read come from?

Are there any differences with the reference?

What is the DNA sequence?

What are the quality scores for each base in the read?

What do we know about the mate?

Which read group does the read belong to?

<http://samtools.github.io/hts-specs/SAMv1.pdf>

Slide courtesy of Andrew Farrell

What needs to be stored?



Where did the read map?

How confident are we that we are correct?

Which strand does the read come from?

Are there any differences with the reference?

What is the DNA sequence?

What are the quality scores for each base in the read?


What do we know about the mate?

Which read group does the read belong to?

<http://samtools.github.io/hts-specs/SAMv1.pdf>

Slide courtesy of Andrew Farrell

What needs to be stored?



Where did the read map?

How confident are we that we are correct? Which strand does the read come from?

Are there any differences with the reference?

What is the DNA sequence?

What are the quality scores for each base in the read?


What do we know about the mate?

Which read group does the read belong to?

<http://samtools.github.io/hts-specs/SAMv1.pdf>

Slide courtesy of Andrew Farrell

Store the alignment



Standardize alignment formats

- SAM : Sequence Alignment/Map
- BAM: Binary Alignment/Map
- Can be indexed allowing fast access of regions
- Simple format
- Can represent single and paired end reads
- Many toolkits now available to process data

<http://samtools.github.io/hts-specs/SAMv1.pdf>

Slide courtesy of Andrew Farrell

SAM/BAM format

Aspect	SAM	BAM
Format	Plain text (human-readable)	Binary (compressed)
File Size	Larger, less efficient	Smaller, more efficient
Access Speed	Slower for large files	Faster, especially with indexing
Readability	Directly readable	Needs tools (e.g., <code>samtools</code>)

- BAM is a compressed version of SAM; lossless BGZF format
- BAM files are usually ‘indexed’
 - A ‘.bai’ file will be found beside the ‘.bam’ file
 - Indexing provides fast retrieval of alignments overlapping a specified region without going through all alignments.
 - BAM must be sorted by the reference ID and then the leftmost coordinate before indexing
 - <http://samtools.sourceforge.net/SAM1.pdf>

Example SAM/BAM/CRAM header section (abbreviated)

Example SAM/BAM/CRAM alignment section (only 10 alignments shown)

29

A SAM/BAM file is divided in header & alignment sections

Example SAM/BAM header section (abbreviated)

```
mgriffit@linus270 ~$ samtools view -H /gscmnt/gc13001/info/model_data/2891632684/build136494552/alignments/136080019.bam | grep -P "SN\|22\|HD\|RG\|PG"
@HD    VN:1.4  SO:coordinate
@SQ    SN:22  LN:51304566  UR:ftp://ftp.ncbi.nih.gov/genbank/genomes/Eukaryotes/vertebrates_mammals/Homo_sapiens/GRCh37/special_requests/GRCh37-lite.fa.gz AS:GRCh37-lite  M5:a718acaa6135fdca8357d5bfe9
4211dd SP:Homo sapiens
@RG    ID:2888721359  PL:illumina  PU:D1BA4ACXX.3  LB:H_KA-452198-0817007-cDNA-3-lib1  PI:365  DS:paired end  DT:2012-10-03T19:00:00-0500  SM:H_KA-452198-0817007  CN:WUGSC
@PG    ID:2888721359  VN:2.0.8  CL:tophat --library-type fr-secondstrand --bowtie-version=2.1.0
@PG    ID:MarkDuplicates  PN:MarkDuplicates  PP:2888721359  VN:1.85(exported)  CL:net.sf.picard.sam.MarkDuplicates INPUT=[/gscmnt/gc13001/info/build_merged_alignments/merged-alignment-blade10-2-5.gsc.wustl.edu-jwalker-15434-136080019/scratch-ILg6Y/H_KA-452198-0817007-cDNA-3-lib1-2888360300.bam] OUTPUT=/gscmnt/gc13001/info/build_merged_alignments/merged-alignment-blade10-2-5.gsc.wustl.edu-jwalker-15434-136080019/scratch-ILg6Y/H_KA-452198-0817007-cDNA-3-lib1-2888360300-post_dup.bam METRICS_FILE=/gscmnt/gc13001/info/build_merged_alignments/merged-alignment-blade10-2-5.gsc.wustl.edu-jwalker-15434-136080019/staging-1iuJS/H_KA-452198-0817007-cDNA-3-lib1-2888360300.metrics REMOVE_DUPLICATES=false ASSUME_SORTED=true MAX_FILE_HANDLES_FOR_READ_ENDS_MAP=9500 TMP_DIR=[/gscmnt/gc13001/info/build_merged_alignments/merged-alignment-blade10-2-5.gsc.wustl.edu-jwalker-15434-136080019/scratch-ILg6Y] VALIDATION_STRINGENCY=SILENT MAX_RECORDS_IN_RAM=500000 PROGRAM_RECORD_ID=MarkDuplicates PROGRAM_GROUP_NAME=MarkDuplicates MAX_SEQUENCES_FOR_DISK_READ_ENDS_MAP=50000 SORTING_COLLECTION_SIZE_RATIO=0.25 READ_NAME_REGEX=[a-zA-Z0-9]+:[0-9]+:[0-9]+:[0-9]+:[0-9]+.* OPTICAL_DUPLICATE_PIXEL_DISTANCE=100 VERBOSITY=INFO QUIET=false COMPRESSION_LEVEL=5 CREATE_INDEX=false CREATE_MD5_FILE=false
mgriffit@linus270 ~$
```

```
@HD    VN:1.3  SO:coordinate
@SQ    SN:20   LN:63025520
@RG    ID:HG00096  SM:HG00096
@PG    ID:HG00096  PN:bwa  CL:/Users/AlistairNWard/Work/gkno/gkno_launcher/tools/bwa/bwa mem -t
```

Version (VN) and sort order (SO) - Important!

Reference sequence (SQ) and sequence length (LN)

Read group (RG) and sample (SM)

Programs (PG) that have been run on the data

Slide courtesy of Andrew Farrell

SAM/BAM header section

- Used to describe source of data, reference sequence, method of alignment, etc.
- Each section begins with character '@' followed by a two-letter record type code. These are followed by two-letter tags and values:
 - @HD The header line
 - VN: format version
 - SO: Sorting order of alignments
 - @SQ Reference sequence dictionary
 - SN: reference sequence name
 - LN: reference sequence length
 - SP: species
 - @RG Read group
 - ID: read group identifier
 - CN: name of sequencing center
 - SM: sample name
 - @PG Program
 - PN: program name
 - VN: program version

Example SAM/BAM alignment section (only 10 alignments shown)

Slide courtesy of Obi and Malachi Griffith

SAM/BAM alignment section

Col	Field	Type	Regex/Range	Brief description
1	QNAME	String	[!-?A-~]{1,255}	Query template NAME
2	FLAG	Int	[0,2 ¹⁶ -1]	bitwise FLAG
3	RNAME	String	* [!-()+-<>-~] [!-~]*	Reference sequence NAME
4	POS	Int	[0,2 ²⁹ -1]	1-based leftmost mapping POSition
5	MAPQ	Int	[0,2 ⁸ -1]	MAPping Quality
6	CIGAR	String	* ([0-9]+[MIDNSHPX=])+	CIGAR string
7	RNEXT	String	* = [!-()+-<>-~] [!-~]*	Ref. name of the mate/next segment
8	PNEXT	Int	[0,2 ²⁹ -1]	Position of the mate/next segment
9	TLEN	Int	[-2 ²⁹ +1,2 ²⁹ -1]	observed Template LENgth
10	SEQ	String	* [A-Za-z=.]+	segment SEQUENCE
11	QUAL	String	[!-~]+	ASCII of Phred-scaled base QUALity+33

SRR062634.14576120	163	20	899919	60	100M	=	900037	218	TTCCCCAGTAGCTGGGATTACAGGCATACGCCACCATC
--------------------	-----	----	--------	----	------	---	--------	-----	--

1	QNAME	e.g. HWI-ST495_129147882:1:2302:10269:12362
2	FLAG	e.g. 163
3	RNAME	e.g. 20
4	POS	e.g. 899919
5	MAPQ	e.g. 60
6	CIGAR	e.g. 100M
7	RNEXT	e.g. =
8	PNEXT	e.g. 900037
9	TLEN	e.g. 218
10	SEQ	e.g. CCTGTTTCTCCACAAAGTGTTTACTTTTGGATTTTGGCAGTCTAACAGGTGAAGCCCTGGAGATTCTTATTAGTGATTTGGGCTGGGGCCTGGCCATGT
11	QUAL	e.g. CCCFFFFFFHHHHHJJJIJFIJJJJJJJJJJJHIIJJJJJJJJJJGHGHIJHIIJJJJJJJJJGHGGIJJJJJJJIJEEHHHHFFFFCDDDDDDDDDB@ACDD

SAM/BAM flags explained

- 12 bitwise flags describing the alignment
- Stored as a binary string of length 12 instead of 12 columns of data
- Value of '1' indicates the flag is set. e.g. 001000000000
- All combinations can be represented as a number from 0 to 4095 (i.e. $2^{12}-1$). This number is used in the BAM/SAM file.
- You can specify 'required' or 'filter' flags in samtools view using the '-f' and '-F' options respectively

Bit		Description
1	0x1	template having multiple segments in sequencing
2	0x2	each segment properly aligned according to the aligner
4	0x4	segment unmapped
8	0x8	next segment in the template unmapped
16	0x10	SEQ being reverse complemented
32	0x20	SEQ of the next segment in the template being reverse complemented
64	0x40	the first segment in the template
128	0x80	the last segment in the template
256	0x100	secondary alignment
512	0x200	not passing filters, such as platform/vendor quality controls
1024	0x400	PCR or optical duplicate
2048	0x800	supplementary alignment

Note that to maximize confusion, each bit is described in the SAM specification using its hexadecimal representation (i.e., '0x10' = 16 and '0x40' = 64).

<http://broadinstitute.github.io/picard/explain-flags.html>


CIGAR strings explained

- The 'CIGAR' (**C**ompact **I**diosyncratic **G**apped **A**lignment **R**eport)
- The CIGAR string is a sequence of base lengths and associated 'operations' indicating which bases align to the reference (either a match or mismatch), are deleted, are inserted, represent introns, etc.

Op	BAM	Description
M	0	alignment match (can be a sequence match or mismatch)
I	1	insertion to the reference
D	2	deletion from the reference
N	3	skipped region from the reference
S	4	soft clipping (clipped sequences present in SEQ)
H	5	hard clipping (clipped sequences NOT present in SEQ)
P	6	padding (silent deletion from padded reference)
=	7	sequence match
X	8	sequence mismatch

Reference: ACTTTTCATCCCTAAA---CAACC---CTGTGTTTCCC

Sample: AC**CGG**TCAT---TAA**TT**CAACC**CTT**CTGTG**AAA**TCCC



CRAM files

- CRAM is an ultra-compressed version of a BAM
 - Usually between 30-60% smaller than the corresponding BAM
- Stores “diffs” from the reference genome
 - requires the matching reference genome to restore original data!
- Base quality binning may be used as well
- Some tools still require conversion back to bam

SAM > BAM > CRAM

Quality Score Bins	Example of Empirically Mapped Quality Scores*
N (no call)	N (no call)
2–9	6
10–19	15
20–24	22
25–29	27
30–34	33
35–39	37
≥ 40	40

By replacing the quality scores between 19 and 25 with a new score of 22, data storage space is conserved.

*The mapped quality score of each bin (except “N”) is subject to change depending on individual Q-tables.

Introduction to the BED format

- When working with BAM files, it is very common to want to examine a focused subset of the reference genome, e.g. the exons of a gene
- These subsets are commonly specified in 'BED' files (Browser Extensible Data)
<https://genome.ucsc.edu/FAQ/FAQformat.html#format1>
- Many BAM manipulation tools accept regions of interest in BED format
- Basic BED format (tab separated):
 - Chromosome name, start position, end position (BED3)
 - Coordinates in BED format are 0 based:
 - Start position (the 0-based start coordinate)
 - End position (the 1-based end coordinate)

Introduction to the BED format

- There are several flavors of BED format: BED3, BED4, BED6, BED8, etc.
- First 3 fields always required: chr, start, stop
- Followed by up to 9 additional optional fields: name, score, strand, thickStart, thickEnd, itemRGB, blockCount, blockSizes, blockStarts

chr7	127471196	127472363	Pos1	0	+
chr7	127472363	127473530	Pos2	0	+
chr7	127473530	127474697	Pos3	0	+
chr7	127474697	127475864	Pos4	0	+
chr7	127475864	127477031	Neg1	0	-
chr7	127477031	127478198	Neg2	0	-
chr7	127478198	127479365	Neg3	0	-
chr7	127479365	127480532	Pos5	0	+
chr7	127480532	127481699	Neg4	0	-

Manipulation of SAM/BAM and BED files


- Several tools are used ubiquitously in sequence analysis to manipulate these files
- SAM/BAM files
 - samtools
 - bamtools
 - Picard
- BED files
 - bedtools
 - bedops



Common sources of confusion

- Genomic coordinate systems
- Genome builds
- Variant representation

Genomic coordinates – 1 vs 0 based



	1-based	0-based
Indicate a single nucleotide	chr1:4-4 G	chr1:3-4 G
Indicate a range of nucleotides	chr1:2-4 ACG	chr1:1-4 ACG
Indicate a single nucleotide variant	chr1:5-5 T/A	chr1:4-5 T/A


- 1-based : Single nucleotides, variant positions, or ranges are specified directly by their corresponding nucleotide numbers
 - GFF, SAM, VCF, Ensembl browser, ...
- 0-based: Single nucleotides, variant positions, or ranges are specified by the coordinates that flank them
 - BED, BAM, UCSC browser, ...

Reference Genome builds

Lift-over

Current human: GRCh38, hg38, b38
alternates: GRCh38v2_ccdg,
GRCh38_full_analysis_set_plus_decoy_hla
Previous human: GRCh37, hg19, b37

Current mouse: mm11
Still used mouse: GRCm38, mm10



For a detailed discussion of various human reference genome flavors refer here:
https://pmbio.org/module-02-inputs/0002/02/01/Reference_Genome/

Variant shifting (alignment) and parsimony/trimming

Reference and alternative alleles of a CA short tandem repeat (STR)

REF
ALT

GGGCACACACAGGG
GGGCACACAGGG

← CA deletion from the reference

Genome Reference		Variant Call Format			
GGGCACACACAGGG		POS	REF	ALT	
REF	CA	8	CA	.	Not left aligned and alternate allele is empty
ALT	.				
REF	CAC	6	CAC	C	Not left aligned but parsimonious
ALT	C				
REF	GCACA	3	GCACA	GCA	Not right trimmed
ALT	GCA				
REF	GGCA	2	GGCA	GG	Not left trimmed
ALT	GG				
REF	GCA	3	GCA	G	Normalized (left aligned & parsimonious)
ALT	G				

Alleles represented against the human genome reference. Allele pairs are colored the same, all are representations of the same variant.

Alleles represented in Variant Call Format, all are representations of the same variant.

Parsimony: representing variant in as few nucleotides as possible without reducing the length of any allele to 0

Left (right) aligning = shifting the start position of a variant as far to the left (right) as possible

How should I sort my SAM/BAM file?

- Generally BAM files are sorted by position
 - This is for performance reasons
 - When sorted and indexed, arbitrary positions in a massive BAM file can be accessed rapidly
- Certain tools require a BAM sorted by read name
 - Usually this is when we need to easily identify both reads of a pair
 - The insert size between two reads may be large
 - In fusion detection we are interested in read pairs that map to different chromosomes

Unaligned

FASTQ (seq + quality)
FASTA(seq reads only)

Aligned

SAM (human readable, everything, HUGE)
BAM (binary, everything)
CRAM (smaller, loss of info)
BED (smallest, text ,coordinates only)



size

We are on a Coffee Break & Networking
Session